

## NovaSpark Technologies - Company Overview

NovaSpark Technologies is a B2B software company founded in 2017 and headquartered in Austin, Texas. We specialise in AI-powered data infrastructure tools that help mid-market and enterprise teams ingest, process, and act on their data in real time. As of 2025 we serve over 1,400 customers across North America, Europe, and the Asia-Pacific region, ranging from 50-person scale-ups to Fortune 500 organisations.

Our engineering team of 280 employees ships weekly releases across five product lines. We are a remote-first company with offices in Austin (HQ), London, and Singapore.

## Mission & Vision

**Mission:** Empower every data team to make confident decisions at the speed of their business.

**Vision:** A world where latency between raw data and business insight is measured in seconds, not days.

Core values:

- \* Transparency - we publish our roadmap publicly and hold monthly town-halls with customers.
- \* Reliability - we target 99.95% uptime across all managed services and compensate customers automatically when SLAs are missed.
- \* Developer-first - every feature ships with a CLI, an API, and a UI so teams can integrate however they prefer.

## Company History

2017 - NovaSpark founded by Elena Okafor and James Tran after four years at a major cloud provider. Seed round of \$3.2 M closed.

2019 - Series A (\$18 M). SparkStream v1 launched, signing 120 customers in the first six months.

2021 - Series B (\$65 M). Acquired DataVault Inc. to add encryption-at-rest and key-management capabilities to the platform.

2023 - Series C (\$140 M). Expanded to EMEA and APAC. Launched SparkML and SparkGovernance. ARR crossed \$42 M.

2025 - 1,400+ customers, 280 employees, ARR of \$89 M.

## Product Catalog

NovaSpark offers five integrated products. Each can be purchased standalone or as part of the NovaSpark Platform bundle.

### 1. SparkStream - Real-Time Data Pipeline

SparkStream is our flagship product: a fully managed, serverless event streaming platform built on Apache Kafka under the hood. It handles ingestion, transformation, and routing of event data at up to 10 million events per second per tenant.

Key features:

- \* 200+ pre-built connectors (Salesforce, Snowflake, PostgreSQL, S3, Webhook, and more).
- \* Sub-100 ms end-to-end latency at p99.
- \* Schema registry with auto-evolution and backward compatibility checks.
- \* Built-in dead-letter queue and automatic retry with exponential backoff.
- \* Exactly-once delivery semantics across all connectors.
- \* Visual pipeline builder (drag-and-drop) and full Terraform provider.

### 2. SparkDB - Analytical Data Warehouse

SparkDB is a columnar, ANSI-SQL analytical database optimised for workloads that mix batch reporting with interactive BI queries. It separates compute and storage, allowing teams to pause clusters and pay only for storage when idle.

Key features:

- \* Petabyte-scale storage on object storage (S3-compatible).
- \* Automatic workload isolation: separate compute clusters per team.
- \* Time-travel queries: restore any table to any point in the last 90 days.
- \* Native integrations with dbt, Apache Spark, and Jupyter.
- \* Row-level security and column masking built in.
- \* Query result cache with configurable TTL.

### 3. SparkML - Machine Learning Platform

SparkML is an end-to-end MLOps platform that takes teams from feature engineering to model serving without leaving the NovaSpark ecosystem.

Key features:

- \* Managed feature store with point-in-time correct joins.
- \* AutoML for classification, regression, and time-series forecasting.
- \* GPU-accelerated training clusters (H100 and A100) with spot pricing.
- \* One-click model deployment as REST endpoint or batch inference job.
- \* Experiment tracking (MLflow-compatible API).
- \* Drift detection and automated re-training triggers.

## 4. SparkGovernance - Data Catalog & Lineage

SparkGovernance gives data teams a single pane of glass over all their data assets: tables, dashboards, ML models, and pipelines.

Key features:

- \* Automated metadata harvesting via passive query log analysis.
- \* Column-level lineage graph rendered as an interactive DAG.
- \* Business glossary with ownership and certification workflows.
- \* PII classification and GDPR / CCPA compliance reporting.
- \* Integration with Jira and Slack for data-quality incident routing.

## 5. SparkAlert - Anomaly Detection & Observability

SparkAlert monitors data pipelines and warehouse queries for anomalies, SLA breaches, and cost spikes, alerting the right person before downstream consumers notice a problem.

Key features:

- \* ML-based anomaly detection with zero manual threshold-setting.
- \* Freshness, volume, schema, and distribution monitors out of the box.
- \* PagerDuty, OpsGenie, Slack, and email integrations.
- \* Cost anomaly detection: flags when a query or pipeline exceeds budget.
- \* Root-cause analysis assistant powered by LLM.

## Pricing Tiers

All plans are billed annually. Monthly billing is available at a 15% premium. All prices are per product unless the Platform bundle is chosen.

**Starter - \$499 / month**

- \* Up to 5 users.
- \* 50 GB storage (SparkDB) or 10 M events/month (SparkStream).
- \* Community support (forum + docs).
- \* 99.9% uptime SLA.
- \* Best for: small teams evaluating NovaSpark in production.

**Professional - \$1,999 / month**

- \* Up to 25 users.
- \* 500 GB storage or 200 M events/month.
- \* Email and chat support (business hours, 4-hour response SLA).
- \* 99.95% uptime SLA.
- \* SSO (SAML 2.0, OIDC).
- \* Best for: growing data teams with production workloads.

**Enterprise - Custom pricing**

- \* Unlimited users.
- \* Custom storage and event-volume commitments.
- \* 24/7 dedicated support with a named Customer Success Manager.
- \* 99.99% uptime SLA with financial penalties for breaches.

- \* Private networking (VPC peering, PrivateLink).
- \* Custom data retention up to 10 years.
- \* HIPAA BAA, SOC 2 Type II, ISO 27001 compliance packages available.
- \* On-premise / hybrid deployment option.
- \* Best for: large enterprises with compliance, scale, or custom needs.

#### Platform Bundle - 30% discount

\* All five products at the same tier for 30% off the sum of individual list prices. Ideal for teams adopting the full NovaSpark ecosystem.

## Security & Compliance

Security is a first-class concern at NovaSpark. We invest heavily in infrastructure hardening, third-party audits, and customer-controlled encryption.

Certifications & standards:

- \* SOC 2 Type II (annual audit, report available on request).
- \* ISO 27001 certified since 2022.
- \* GDPR and CCPA compliant. DPA available on the Enterprise plan.
- \* HIPAA Business Associate Agreement available for Enterprise customers.

Encryption:

- \* All data encrypted at rest with AES-256.
- \* All data in transit encrypted with TLS 1.3.
- \* Customer-managed encryption keys (CMEK) available on Enterprise.
- \* Key rotation on a configurable schedule (default: 90 days).

Access control:

- \* Role-based access control (RBAC) with 12 built-in roles.
- \* Attribute-based access control (ABAC) on Enterprise.
- \* Multi-factor authentication enforced for all accounts.
- \* SSO via SAML 2.0 and OIDC on Professional and Enterprise plans.
- \* Audit logs retained for 12 months (36 months on Enterprise).

Network security:

- \* IP allowlisting available on all paid plans.
- \* VPC peering and AWS PrivateLink on Enterprise.
- \* Quarterly penetration tests by independent security firms.
- \* Bug bounty programme managed via HackerOne.

## Support & Contact Information

Support channels (vary by plan):

- \* Documentation portal: [docs.novaspark.io](https://docs.novaspark.io)
- \* Community forum: [community.novaspark.io](https://community.novaspark.io)
- \* In-app chat (Professional & Enterprise): business hours, 4-hour first response.
- \* Email: [support@novaspark.io](mailto:support@novaspark.io)
- \* Dedicated phone line (Enterprise only): +1-800-NOVA-SPK
- \* Slack Connect channel (Enterprise only): direct line to your CSM.

Sales enquiries:

- \* Email: [sales@novaspark.io](mailto:sales@novaspark.io)
- \* Book a demo: [novaspark.io/demo](https://novaspark.io/demo)
- \* Partner programme: [partners@novaspark.io](mailto:partners@novaspark.io)

Headquarters:

NovaSpark Technologies Inc.  
500 Congress Ave, Suite 2200

## Frequently Asked Questions

Q: Can I try NovaSpark before purchasing?

A: Yes. Every product offers a 14-day free trial with full functionality and no credit card required. Trials can be extended to 30 days on request.

Q: Do you support on-premise deployment?

A: On-premise and hybrid (cloud + on-prem) deployments are available exclusively on the Enterprise plan. We provide a Kubernetes Helm chart and dedicated implementation support.

Q: How is data billed in SparkDB?

A: Storage is billed per GB-month of compressed data. Compute is billed per vCPU-hour consumed by active query clusters. Paused clusters incur zero compute costs.

Q: What happens if I exceed my event quota in SparkStream?

A: Events over the monthly quota are metered at \$0.12 per million additional events. You receive an email alert at 80% and 100% of your quota.

Q: Is there a Service Level Agreement for uptime?

A: Yes. Starter: 99.9% (allows ~8.7 h downtime/year). Professional: 99.95% (~4.4 h). Enterprise: 99.99% (~52 min). Enterprise customers receive automatic service credits for any breach.

Q: Can I export my data if I cancel?

A: Absolutely. You can export all data in open formats (Parquet, CSV, JSON) at any time. After cancellation we retain your data for 30 days before permanent deletion, giving you time to complete the export.